

NON-WSQ COURSE

Agentic AI for Video Creation

Hermes → create → review → approve → publish

C436 | v2.0 | 2 days / 15 hours

Tertiary Infotech Academy Pte Ltd
UEN 201200696W

MiniMax M3 • Eight self-contained Hermes labs



Trainer profile



Your trainer facilitates evidence-led AI workflow design and video-production governance.

1

Hermes Agent architecture

2

Generative media production controls

3

Video tools and custom skills

4

Technical assessment and remediation

Dr Alfred Ang



Principal Trainer, Tertiary Infotech Academy

1

Doctorate-level applied technology educator

2

AI, data, cloud and automation curriculum leadership

3

Industry-oriented assessment and implementation practice

4

Focus: explainable, governed technical outcomes

Learner introduction

YOU

Connect your existing role to one video-production automation problem.

1

Role and organisation context

2

Current video or content workflow

3

One repeated production bottleneck

4

One risk you want the automation to control

Working agreements



RULE

Use governed experimentation throughout the course.

1

Use synthetic or approved data only

2

Never paste live secrets into prompts or lab files

3

Keep YouTube uploads private during labs

4

Challenge outputs with evidence and timecodes

Course facts and competency context



Non-WSQ commercial short course

1 Course code: C436

2 Duration: 2 days / 15 instructional hours

3 Assessment: none

4 Version: v2.0 / 16 September 2026

Day 1 lesson plan

D1

Hermes setup, MiniMax M3 and governed video prompt engineering.

1

Hermes Desktop installation and diagnostics

2

MiniMax-M3 provider configuration

3

FRAME-CUT prompt contracts

4

Labs 01–03: setup to simple video

Day 2 lesson plan

D2

Video tools, skills, multi-agent control and scheduled release.

1

Remotion, Manim, Higgsfield and FFmpeg

2

Custom branding, tone and style skill

3

Kanban review and private YouTube upload

4

Labs 04–08 plus 2-hour assessment

Learning outcomes

LO

At course completion, you can design, assess and remediate a governed agentic-video pipeline.

1

LO1: Develop an editing strategy and work plan in Hermes Agent using MiniMax M3, structured video prompts and governed tool selection.

2

LO2: Create and customise an evidence-backed video with Hermes skills, Remotion, Manim, Higgsfield-compatible requests and deterministic media checks.

3

LO3: Orchestrate research, production, review and YouTube release agents through a durable Kanban board and controlled scheduled publishing.

Access the labs

LAB

Eight self-contained lab folders carry the complete Hermes production workflow.

1

Open each folder as a Hermes project

2

Use only supplied synthetic inputs

3

Use AI-PROMPTS.md or its PDF twin

4

Run verify.py before moving on

<https://github.com/tertiarycourses/C436-Agentic-AI-for-Video-Creation>

1

TOPIC 1

Hermes Agent Setup, MiniMax M3 and Video Prompt Engineering

LO1 - configure Hermes safely and convert creative intent into a machine-checkable video plan

1

CONTROL

2

EVIDENCE

3


AUTOMATION

4


DECISION

Every mechanism ends in observable evidence—not an agent's assertion.

01 · Hermes Desktop runtime boundary: execution path



01 · Hermes Desktop runtime boundary: contract and control



Use the official desktop page and stop when hermes doctor reports a failed dependency.

01 · Hermes Desktop runtime boundary: evidence and decision

MEASURE

setup pass = required diagnostics passed / required diagnostics

ACCEPTANCE EVIDENCE

setup-evidence.json records official source, version and diagnostic status.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

02 · Shared desktop and CLI state: execution path



02 · Shared desktop and CLI state: contract and control



Name the active profile and compare provider, model and skill directory before production.

02 · Shared desktop and CLI state: evidence and decision

MEASURE

**state consistency = matching settings /
checked settings**

ACCEPTANCE EVIDENCE

profile-check.json shows the same
provider and model on both surfaces.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

03 · MiniMax M3 provider contract: execution path



03 · MiniMax M3 provider contract: contract and control



Use provider minimax with model MiniMax-M3 and verify the returned model identifier.

03 · MiniMax M3 provider contract: evidence and decision

MEASURE

model handshake pass = successful tool-capable response

ACCEPTANCE EVIDENCE

model-check.json records MiniMax-M3 without recording the key.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

04 · Trial and quota verification: execution path



04 · Trial and quota verification: contract and control



Treat the offer as time-sensitive; record live terms at sign-up and prepare a fallback.

04 · Trial and quota verification: evidence and decision

MEASURE

budget headroom = available quota - reserved quota

ACCEPTANCE EVIDENCE

trial-checklist.md contains observed terms, date and learner decision.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER

named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

05 · Credential isolation: execution path




STATE PASSED BETWEEN STAGES

secret_name	scope
provider	rotation_due

FAILURE MODE

A MiniMax or YouTube key is pasted into a prompt, screenshot or repository.

05 · Credential isolation: contract and control



Keep secrets in Hermes credential/config storage and publish placeholders only.

05 · Credential isolation: evidence and decision

MEASURE

secret exposure count = 0

ACCEPTANCE EVIDENCE

secret scan reports zero live tokens in the lab package.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

06 · Model capability routing: execution path



06 · Model capability routing: contract and control



Use M3 for reasoning and tool orchestration; use a video renderer or generator for media.

06 · Model capability routing: evidence and decision

MEASURE

routing yield = accepted outputs / routed calls

ACCEPTANCE EVIDENCE

tool-routing.csv links every task to a capable model or tool.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

07 · FRAME-CUT prompt contract: execution path



07 · FRAME-CUT prompt contract: contract and control



Use FRAME-CUT and validate the returned shot plan before generation.

07 · FRAME-CUT prompt contract: evidence and decision

MEASURE

prompt validity = required fields present / required fields

ACCEPTANCE EVIDENCE

prompt-contract.json passes schema and duration checks.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

08 · Format and finish: execution path



08 · Format and finish: contract and control



Freeze delivery properties before shot generation and assembly.

08 · Format and finish: evidence and decision

MEASURE

delivery compliance = passed technical fields / required fields

ACCEPTANCE EVIDENCE

video-spec.json contains approved format and finish values.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

09 · Role, references and assets: execution path



09 · Role, references and assets: contract and control



Limit reasoning to named references and a manifest of authorized assets.

09 · Role, references and assets: evidence and decision

MEASURE

asset coverage = registered assets / required assets

ACCEPTANCE EVIDENCE

asset-manifest.csv resolves every asset and source identifier.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

10 · Action, motion and environment: execution path



10 · Action, motion and environment: contract and control



Use one dominant subject action and one camera move per timed shot.

10 · Action, motion and environment: evidence and decision

MEASURE

usable shot rate = accepted shots / attempts

ACCEPTANCE EVIDENCE

shot-list.json shows one action-motion pair per shot.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

11 · Continuity and negative constraints: execution path



11 · Continuity and negative constraints: contract and control



Reuse continuity tokens and compare keyframes against explicit invariants.

11 · Continuity and negative constraints: evidence and decision

MEASURE

continuity defects per reviewed shot

ACCEPTANCE EVIDENCE

continuity-review.json records deviations and acceptance.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

12 · Structured shot-plan output: execution path



12 · Structured shot-plan output: contract and control



Require strict JSON and fail on gaps, overlaps or unknown fields.

12 · Structured shot-plan output: evidence and decision

MEASURE

timeline variance = planned duration - target duration

ACCEPTANCE EVIDENCE

shot-plan.json totals the target duration with no overlap.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

13 · Simple video agent loop: execution path



13 · Simple video agent loop: contract and control



Set one render attempt, one review pass and an explicit repair decision.

13 · Simple video agent loop: evidence and decision

MEASURE

first-pass completion rate

ACCEPTANCE EVIDENCE

run-ledger.json contains the plan, call, result and stop reason.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION

PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

14 · Cost, latency and stop rules: execution path



14 · Cost, latency and stop rules: contract and control



Cap attempts and require a human decision before additional paid generation.

14 · Cost, latency and stop rules: evidence and decision

MEASURE

budget variance = actual cost - approved cost

ACCEPTANCE EVIDENCE

budget-ledger.csv reconciles attempts, latency and cost.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

Boundary map: agent reasoning versus deterministic control




State model: READY, BLOCKED, REPAIR and APPROVED




Only APPROVED can cross a consequential publishing boundary.

Evidence chain: source → asset → approval → external id




A release is explainable when every identifier resolves to immutable evidence.

Failure taxonomy: transient, contract, policy and human-decision



Evidence coverage closes the release gap




INSIGHT

Approval is not the final checkbox; it is the last link in an evidence chain.

Illustrative synthetic values for mechanism explanation.

Budget control contains retry amplification

SGD per run




INSIGHT

Generation dominates unit cost, so retries need deterministic stop rules.

Illustrative synthetic values for mechanism explanation.

Stage latency reveals the optimisation target



INSIGHT

Optimise the longest stage only after quality and safety gates remain intact.

Illustrative synthetic values for mechanism explanation.

Lab 01 · Set Up Hermes Desktop and Connect MiniMax M3

SCENARIO

Install Hermes Desktop from the official source, configure the MiniMax provider for MiniMax-M3 and prove a tool-capable model handshake without exposing a credential.


EXPECTED OUTPUT

**setup-evidence.json and
redacted diagnostic screenshot**

**Detailed click paths, prompts, troubleshooting and acceptance checks are in the
Learner Guide and lab PDF.**

Self-contained folder: labs/lab-01-setup-hermes-and-connect-minimax-m3/

Lab 01 · Hermes execution contract



Each node consumes named keys, emits a status, and leaves retrievable execution evidence.

Lab 01 · evidence and failure gates

PASS EVIDENCE

Hermes launches; the configured provider is minimax; the model is MiniMax-M3; hermes doctor has no blocking failure; published evidence contains placeholders only.

FAIL CLOSED WHEN

credential appears in JSON; evidence is missing; approval is absent; or the same idempotency key already committed.

AI-PROMPTS.md

copy-ready prompt

data/

synthetic input

AI-PROMPTS.pdf

printable prompts

verify.py

acceptance check

Lab 02 · Prompt Hermes to Create a Simple Video

SCENARIO

Use a copy-ready Hermes prompt to convert a supplied 15-second brief into a validated shot plan and deterministic preview video.


EXPECTED OUTPUT

simple-video.mp4, shot-plan.json and ffprobe.json

Detailed click paths, prompts, troubleshooting and acceptance checks are in the Learner Guide and lab PDF.

Self-contained folder: labs/lab-02-prompt-hermes-to-create-a-simple-video/

Lab 02 · Hermes execution contract



Each node consumes named keys, emits a status, and leaves retrievable execution evidence.

Lab 02 · evidence and failure gates

PASS EVIDENCE

The shot plan is valid JSON, totals 15 seconds, uses only supplied assets, and the generated MP4 passes dimensions, codec and duration checks.

FAIL CLOSED WHEN

credential appears in JSON; evidence is missing; approval is absent; or the same idempotency key already committed.

AI-PROMPTS.md

copy-ready prompt

data/

synthetic input

AI-PROMPTS.pdf

printable prompts

verify.py

acceptance check

Lab 03 · Engineer Video Prompts with FRAME-CUT

SCENARIO

Transform a vague creative request into portable shot prompts using the FRAME-CUT framework and evaluate them with a deterministic rubric.


EXPECTED OUTPUT

prompt-pack.json and prompt-score.csv

Detailed click paths, prompts, troubleshooting and acceptance checks are in the Learner Guide and lab PDF.

Self-contained folder: labs/lab-03-engineer-video-prompts-with-frame-cut/

Lab 03 · Hermes execution contract



Each node consumes named keys, emits a status, and leaves retrievable execution evidence.

Lab 03 · evidence and failure gates

PASS EVIDENCE

Every shot has all FRAME-CUT fields, one dominant action, one camera move, explicit continuity and negative constraints, and a measurable output check.

FAIL CLOSED WHEN

credential appears in JSON; evidence is missing; approval is absent; or the same idempotency key already committed.

AI-PROMPTS.md

copy-ready prompt

data/

synthetic input

AI-PROMPTS.pdf

printable prompts

verify.py

acceptance check

Topic 1 technical checkpoint

1

Which state field prevents a silent downstream assumption?

2

What evidence proves the mechanism completed successfully?

3

Which failure must stop rather than retry?

4

Who owns the consequential decision?

Topic 1 transfer challenge

1

Change one input constraint and predict the affected nodes.

2

Name the identifier that preserves provenance across the change.

3

Define one machine threshold and one human decision.

4

Describe the minimum evidence required to resume the run.

2

TOPIC 2

Video Tools, Hermes Skills and Custom Brand Production

LO2 - create a technically valid branded video by combining agent reasoning with specialist media tools

1

CONTROL

2

EVIDENCE

3


AUTOMATION

4


DECISION

Every mechanism ends in observable evidence—not an agent's assertion.

01 · Tool versus skill decision: execution path



01 · Tool versus skill decision: contract and control



Use a tool for precise integrated execution and a skill for repeatable instructions plus scripts.

01 · Tool versus skill decision: evidence and decision

MEASURE

integration fit = satisfied requirements / requirements

ACCEPTANCE EVIDENCE

decision-record.md states why each capability is a tool or skill.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER

named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

02 · Hermes skill anatomy: execution path




STATE PASSED BETWEEN STAGES	
name	description
version	skill_dir

FAILURE MODE

A vague description prevents Hermes from discovering the skill.

02 · Hermes skill anatomy: contract and control



Describe the exact trigger and include deterministic verification commands.

02 · Hermes skill anatomy: evidence and decision

MEASURE

skill test pass = expected artifacts produced

ACCEPTANCE EVIDENCE

SKILL.md and test transcript show correct activation.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER

named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

03 · Skill discovery and progressive loading: execution path




STATE PASSED BETWEEN STAGES

skill_name	category
file_path	task

FAILURE MODE

The agent loads a large library or the wrong similarly named skill.

03 · Skill discovery and progressive loading: contract and control



Search first, view the selected skill, then load only referenced resources.

03 · Skill discovery and progressive loading: evidence and decision

MEASURE

context efficiency = loaded relevant files / loaded files

ACCEPTANCE EVIDENCE

skill-usage.json records the selected skill and resources.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER

named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

04 · Remotion composition contract: execution path




STATE PASSED BETWEEN STAGES

composition_id	fps
duration_frames	input_props

FAILURE MODE

The render duration or aspect ratio drifts from the approved video specification.

04 · Remotion composition contract: contract and control



Derive frames from duration and validate width, height, fps and codec after render.

04 · Remotion composition contract: evidence and decision

MEASURE

duration seconds = duration frames / fps

ACCEPTANCE EVIDENCE

remotion-render.json and ffprobe.json agree with the contract.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER

named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

05 · Manim explanatory scene: execution path



STATE PASSED BETWEEN STAGES

scene_class

frame_rate


resolution

output_file

FAILURE MODE

A mathematical or process animation is visually attractive but semantically wrong.

05 · Manim explanatory scene: contract and control



Bind labels and values to the approved evidence model and review rendered frames.

05 · Manim explanatory scene: evidence and decision

MEASURE

animation timing variance

ACCEPTANCE EVIDENCE

manim-scene.py and contact sheet
prove the intended transformation.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

06 · Higgsfield shot request: execution path



06 · Higgsfield shot request: contract and control



Create a request preview first; call the service only with approved assets and quota.

06 · Higgsfield shot request: evidence and decision

MEASURE

**accepted clip yield = approved clips /
generated clips**

ACCEPTANCE EVIDENCE

higgsfield-request.json records prompt, references and review status.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

07 · FFmpeg deterministic assembly: execution path



07 · FFmpeg deterministic assembly: contract and control



Normalize every source before concat and probe the final container.

07 · FFmpeg deterministic assembly: evidence and decision

MEASURE

render success = valid masters / render attempts

ACCEPTANCE EVIDENCE

ffprobe.json confirms H.264, AAC, dimensions, fps and duration.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

08 · Tool permission boundary: execution path



08 · Tool permission boundary: contract and control



Scope paths and accounts; separate preview from consequential execution.

08 · Tool permission boundary: evidence and decision

MEASURE

unauthorized side effects = 0

ACCEPTANCE EVIDENCE

tool-policy.yaml and audit log show bounded operations.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

09 · Asset provenance and checksum: execution path



09 · Asset provenance and checksum: contract and control



Bind approval to cryptographic hashes and immutable versions.

09 · Asset provenance and checksum: evidence and decision

MEASURE

manifest integrity = verified hashes / manifest rows

ACCEPTANCE EVIDENCE

asset-manifest.csv verifies rights and hashes.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

10 · Brand token system: execution path



10 · Brand token system: contract and control



Store editable brand tokens and reference them from the video skill.

10 · Brand token system: evidence and decision

MEASURE

brand compliance = passed checks / applicable checks

ACCEPTANCE EVIDENCE

brand-profile.yaml and frame samples pass the checklist.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

11 · Tone and style controls: execution path



11 · Tone and style controls: contract and control



Use descriptive attributes and exclusions, not living-artist imitation.

11 · Tone and style controls: evidence and decision

MEASURE

tone adherence score

ACCEPTANCE EVIDENCE

review rubric records tone, pacing and originality evidence.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER

named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

12 · Custom brand-video skill: execution path




STATE PASSED BETWEEN STAGES

skill_name	script_path
template_path	output_contract

FAILURE MODE

The skill depends on hidden context or absolute project paths.

12 · Custom brand-video skill: contract and control



Use Hermes template variables and explicit inputs, outputs and verification.

12 · Custom brand-video skill: evidence and decision

MEASURE

repeatability = matching outputs / repeated runs

ACCEPTANCE EVIDENCE

custom skill produces the expected video and evidence on a clean run.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

13 · Technical video quality gate: execution path



13 · Technical video quality gate: contract and control



Fail closed on missing streams, wrong bounds or unreadable captions.

13 · Technical video quality gate: evidence and decision

MEASURE

technical checks passed / technical checks

ACCEPTANCE EVIDENCE

technical-qc.json and contact sheet show observed values.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

14 · Repair and versioned evidence: execution path



14 · Repair and versioned evidence: contract and control



Write a new immutable version and rerun every affected check.

14 · Repair and versioned evidence: evidence and decision

MEASURE

repair effectiveness = closed findings / reopened findings

ACCEPTANCE EVIDENCE

repair-log.json links finding, change, hash and re-test.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION

PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

Boundary map: agent reasoning versus deterministic control




State model: READY, BLOCKED, REPAIR and APPROVED




Only APPROVED can cross a consequential publishing boundary.

Evidence chain: source → asset → approval → external id



A release is explainable when every identifier resolves to immutable evidence.

Failure taxonomy: transient, contract, policy and human-decision



Evidence coverage closes the release gap




INSIGHT

Approval is not the final checkbox; it is the last link in an evidence chain.

Illustrative synthetic values for mechanism explanation.

Budget control contains retry amplification

SGD per run




INSIGHT

Generation dominates unit cost, so retries need deterministic stop rules.

Illustrative synthetic values for mechanism explanation.

Stage latency reveals the optimisation target



INSIGHT

Optimise the longest stage only after quality and safety gates remain intact.

Illustrative synthetic values for mechanism explanation.

Lab 04 · Install Video Tools and Hermes Skills

SCENARIO

Create a governed tool registry and install learner-safe Hermes skills for Remotion, Manim, Higgsfield request preparation and FFmpeg verification.


EXPECTED OUTPUT

tool-routing.json and skill-smoke-test.json

Detailed click paths, prompts, troubleshooting and acceptance checks are in the Learner Guide and lab PDF.

Self-contained folder: labs/lab-04-install-video-tools-and-skills/

Lab 04 · Hermes execution contract



Each node consumes named keys, emits a status, and leaves retrievable execution evidence.

Lab 04 · evidence and failure gates

PASS EVIDENCE

The registry records capability, installation state, auth mode, side effects and fallback; every supplied SKILL.md passes frontmatter and path checks.

FAIL CLOSED WHEN

credential appears in JSON; evidence is missing; approval is absent; or the same idempotency key already committed.

AI-PROMPTS.md

copy-ready prompt

data/

synthetic input

AI-PROMPTS.pdf

printable prompts

verify.py

acceptance check

Lab 05 · Create a Custom Branded Video Skill

SCENARIO

Turn brand, tone and style rules into a reusable Hermes skill, render a custom video and verify the output against brand and technical evidence.


EXPECTED OUTPUT

custom-video.mp4, brand-review.json and render-evidence.json

Detailed click paths, prompts, troubleshooting and acceptance checks are in the Learner Guide and lab PDF.

Self-contained folder: labs/lab-05-create-custom-branded-video-skill/

Lab 05 · Hermes execution contract



Each node consumes named keys, emits a status, and leaves retrievable execution evidence.

Lab 05 · evidence and failure gates

PASS EVIDENCE

The custom skill is discoverable, uses relative or Hermes template paths, creates an MP4, passes the technical probe and meets every required brand token.

FAIL CLOSED WHEN

credential appears in JSON; evidence is missing; approval is absent; or the same idempotency key already committed.

AI-PROMPTS.md

copy-ready prompt

data/

synthetic input

AI-PROMPTS.pdf

printable prompts

verify.py

acceptance check

Topic 2 technical checkpoint

1

Which state field prevents a silent downstream assumption?

2

What evidence proves the mechanism completed successfully?

3

Which failure must stop rather than retry?

4

Who owns the consequential decision?

Topic 2 transfer challenge

1

Change one input constraint and predict the affected nodes.

2

Name the identifier that preserves provenance across the change.

3

Define one machine threshold and one human decision.

4

Describe the minimum evidence required to resume the run.

3

TOPIC 3

Multi-Agent Kanban, YouTube Release and Scheduled Publishing

LO3 - operate a durable role-based production system with review, release and scheduling safeguards

1

CONTROL

2

EVIDENCE

3


AUTOMATION

4

DECISION

Every mechanism ends in observable evidence—not an agent's assertion.

01 · Agent role contract: execution path




STATE PASSED BETWEEN STAGES

agent_name	allowed_tools
deliverable	done_when

FAILURE MODE

A role such as video agent is too broad to evaluate or recover.

01 · Agent role contract: contract and control



Give every agent one owned deliverable and one measurable completion rule.

01 · Agent role contract: evidence and decision

MEASURE

role completeness = required fields / required fields

ACCEPTANCE EVIDENCE

agent-contracts.yaml validates all four specialist roles.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER

named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

02 · Research agent: execution path




STATE PASSED BETWEEN STAGES

research_task	source_ids
claim_ids	handoff_status

FAILURE MODE

Research output contains fabricated or irretrievable citations.

02 · Research agent: contract and control



Require retrievable URLs, bounded excerpts and explicit limitations.

02 · Research agent: evidence and decision

MEASURE

claim coverage = sourced claims / factual claims

ACCEPTANCE EVIDENCE

research-handoff.json contains approved sources and claim IDs.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER

named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

03 · Video production agent: execution path




STATE PASSED BETWEEN STAGES

claim_ids	shot_plan
skill_names	master_hash

FAILURE MODE

The production agent changes approved claims or bypasses brand controls.

03 · Video production agent: contract and control



Provide immutable research input and require the custom brand-video skill.

03 · Video production agent: evidence and decision

MEASURE

production acceptance = passed render checks / checks

ACCEPTANCE EVIDENCE

video-handoff.json references approved claims and master hash.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

04 · Independent review agent: execution path



04 · Independent review agent: contract and control



Use an independent context and forbid approval authority in the reviewer role.

04 · Independent review agent: evidence and decision

MEASURE

weighted review score and unresolved high count

ACCEPTANCE EVIDENCE

review-handoff.json contains timecoded findings and QA status.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

05 · Upload agent: execution path



05 · Upload agent: contract and control



Default to private and verify the current master hash immediately before upload.

05 · Upload agent: evidence and decision

MEASURE

upload reconciliation = one video ID per idempotency key

ACCEPTANCE EVIDENCE

upload-receipt.json records privacy, disclosure and returned video ID.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

06 · Delegation context contract: execution path



06 · Delegation context contract: contract and control



Include exact paths, inputs, constraints and tests in every delegation context.

06 · Delegation context contract: evidence and decision

MEASURE

handoff pass rate = accepted handoffs / handoffs

ACCEPTANCE EVIDENCE

delegation-log.json records prompt, agent and verified result.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

07 · Parallel work and join: execution path



07 · Parallel work and join: contract and control



Parallelize only independent work and join on explicit verified artifacts.

07 · Parallel work and join: evidence and decision

MEASURE

elapsed reduction with zero skipped dependencies

ACCEPTANCE EVIDENCE

join-ledger.json shows all required parents complete.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER

named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

08 · Durable Hermes Kanban: execution path



STATE PASSED BETWEEN STAGES

task_id

status


assignee

parent_ids

FAILURE MODE

An in-process agent swarm loses state when a session ends.

08 · Durable Hermes Kanban: contract and control



Use Kanban for cross-agent work that must survive restarts and human pauses.

08 · Durable Hermes Kanban: evidence and decision

MEASURE

flow time = done_at - created_at

ACCEPTANCE EVIDENCE

kanban-export.json shows the full dependency chain and statuses.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER

named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

09 · Kanban status and circuit breaker: execution path




STATE PASSED BETWEEN STAGES

status	attempt
failure_class	reviewer

FAILURE MODE

A failing task retries indefinitely or is marked done from a commit alone.

09 · Kanban status and circuit breaker: contract and control



Cap attempts, require contract verification and block repeated failures for intervention.

09 · Kanban status and circuit breaker: evidence and decision

MEASURE

retry recovery and repeated-failure count

ACCEPTANCE EVIDENCE

task history shows checkpoints, review and final evidence.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

10 · Human approval with payload hash: execution path



10 · Human approval with payload hash: contract and control



Bind the named decision to the immutable release package hash.

10 · Human approval with payload hash: evidence and decision

MEASURE

approval integrity = approved hash equals current hash

ACCEPTANCE EVIDENCE

approval-ledger.json proves reviewer, time, decision and hash.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER

named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

11 · YouTube videos.insert contract: execution path



STATE PASSED BETWEEN STAGES

snippet.title

status.privacyStatus


status.containsSyntheticMedia

media_path

FAILURE MODE

The wrong account, privacy state or disclosure reaches YouTube.

11 · YouTube videos.insert contract: contract and control



Confirm channel identity, default private and include applicable synthetic-media disclosure.

11 · YouTube videos.insert contract: evidence and decision

MEASURE

metadata completeness and processing success

ACCEPTANCE EVIDENCE

request-preview.json and private upload receipt match the approved package.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

12 · Publishing idempotency: execution path



12 · Publishing idempotency: contract and control



Reconcile the prior key and video ID before any retry.

12 · Publishing idempotency: evidence and decision

MEASURE

duplicate publish rate = 0

ACCEPTANCE EVIDENCE

publication-ledger.csv contains one committed ID per key.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

13 · Hermes cron job contract: execution path



13 · Hermes cron job contract: contract and control



Put all required paths, gates and stop rules in the job prompt; attach skills explicitly.

13 · Hermes cron job contract: evidence and decision

MEASURE

schedule validity and successful dry run

ACCEPTANCE EVIDENCE

cron-preview.json records schedule, next run and paused state.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION


PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

14 · Scheduled release operations: execution path



14 · Scheduled release operations: contract and control



Schedule preparation and private upload only; require approval before public visibility.

14 · Scheduled release operations: evidence and decision

MEASURE

on-time success rate with zero unauthorized posts

ACCEPTANCE EVIDENCE

operations-ledger.csv records each run, outcome and owner decision.

OBSERVED

machine value

THRESHOLD

approved rule

DECISION

PASS / REPAIR

OWNER


named human

Release principle: evidence travels with the run_id, asset hash and decision—not in a separate chat transcript.

Boundary map: agent reasoning versus deterministic control




State model: READY, BLOCKED, REPAIR and APPROVED



Only APPROVED can cross a consequential publishing boundary.

Evidence chain: source → asset → approval → external id




A release is explainable when every identifier resolves to immutable evidence.

Failure taxonomy: transient, contract, policy and human-decision



Evidence coverage closes the release gap




INSIGHT

Approval is not the final checkbox; it is the last link in an evidence chain.

Illustrative synthetic values for mechanism explanation.

Budget control contains retry amplification

SGD per run




INSIGHT

Generation dominates unit cost, so retries need deterministic stop rules.

Illustrative synthetic values for mechanism explanation.

Stage latency reveals the optimisation target




INSIGHT

Optimise the longest stage only after quality and safety gates remain intact.

Illustrative synthetic values for mechanism explanation.

QA score rises after targeted repair



INSIGHT

Repair should target the failed criterion and trigger a scoped re-test.

Illustrative synthetic values for mechanism explanation.

Lab 06 · Build the Multi-Agent Video Workflow

SCENARIO

Define and simulate four isolated Hermes roles for research, video creation, independent review and approved YouTube upload.


EXPECTED OUTPUT

multi-agent-plan.json and four verified handoff records

Detailed click paths, prompts, troubleshooting and acceptance checks are in the Learner Guide and lab PDF.

Self-contained folder: labs/lab-06-build-multi-agent-video-workflow/

Lab 06 · Hermes execution contract



Each node consumes named keys, emits a status, and leaves retrievable execution evidence.

Lab 06 · evidence and failure gates

PASS EVIDENCE

All roles have bounded tools and outputs; the reviewer is independent; upload is blocked until all parent evidence and the current approval hash pass.

FAIL CLOSED WHEN

credential appears in JSON; evidence is missing; approval is absent; or the same idempotency key already committed.

AI-PROMPTS.md

copy-ready prompt

data/

synthetic input

AI-PROMPTS.pdf

printable prompts

verify.py

acceptance check

Lab 07 · Orchestrate Kanban Review and YouTube Upload

SCENARIO

Create a durable Hermes Kanban dependency chain, review the approved video, and prepare or execute a private YouTube upload with explicit human authorization.


EXPECTED OUTPUT

kanban-export.json, approval-ledger.json and private-upload receipt or dry-run preview

Detailed click paths, prompts, troubleshooting and acceptance checks are in the Learner Guide and lab PDF.

Self-contained folder: labs/lab-07-orchestrate-kanban-review-and-youtube-upload/

Lab 07 · Hermes execution contract



Approval resumes only when the payload hash still matches.



Lab 07 · evidence and failure gates

PASS EVIDENCE

The Kanban graph has no missing dependency; review is required; the request points to the approved master; privacy is private; no credential is present in artifacts.

FAIL CLOSED WHEN

credential appears in JSON; evidence is missing; approval is absent; or the same idempotency key already committed.

AI-PROMPTS.md

copy-ready prompt

data/

synthetic input

AI-PROMPTS.pdf

printable prompts

verify.py

acceptance check

Lab 08 · Schedule Controlled Video Publishing with Hermes Cron

SCENARIO

Create a paused Hermes cron job that prepares a release from a self-contained prompt, dry-run it, and enable a cadence only after evidence review.


EXPECTED OUTPUT

**cron-preview.json and
operations-ledger.csv**

**Detailed click paths, prompts, troubleshooting and acceptance checks are in the
Learner Guide and lab PDF.**

Self-contained folder: labs/lab-08-schedule-controlled-video-publishing/

Lab 08 · Hermes execution contract



Each node consumes named keys, emits a status, and leaves retrievable execution evidence.

Lab 08 · evidence and failure gates

PASS EVIDENCE

The schedule and timezone are explicit, the prompt is self-contained, the custom skill is attached, the job begins paused, and the dry run cannot publish an unapproved or duplicate video.

FAIL CLOSED WHEN

credential appears in JSON; evidence is missing; approval is absent; or the same idempotency key already committed.

AI-PROMPTS.md

copy-ready prompt

data/

synthetic input

AI-PROMPTS.pdf

printable prompts

verify.py

acceptance check

Course achievement map

1

LO1: Develop an editing strategy and work plan in Hermes Agent using MiniMax M3, structured video prompts and governed tool selection.

2

LO2: Create and customise an evidence-backed video with Hermes skills, Remotion, Manim, Higgsfield-compatible requests and deterministic media checks.

3

LO3: Orchestrate research, production, review and YouTube release agents through a durable Kanban board and controlled scheduled publishing.

Release controls that must remain intact

GATE

Automation is complete only when governance travels with the media.

1

Every factual claim resolves to source evidence

2

Every reviewed asset resolves to an immutable hash

3


Every external release has a named human approval

4

Every platform result is reconciled by idempotency key

Build video systems that can explain what happened.

Evidence before release • Human approval before consequence



Dr Alfred Ang | Tertiary Infotech Academy Pte Ltd

C436 • v2.0 • 15 September 2026